STRATEGY BRIEF

AI Adoption Strategy

PREPARED FOR	VERSION	DATE	HORIZON
The Chairman	v1 — Draft	June 2026	24 months

Contents

SECTION 01

Opening Hypothesis

Alfa's Next Identity

Alfa Electronics has, since 1987, been an electronics-hardware company with software embedded inside the boxes it ships. Over the next 24 months, that order inverts. **Alfa becomes an AI-enabled software-on-hardware company** — a company whose products are intelligent software, delivered through the hardware substrates Alfa has spent thirty-nine years learning how to design, install, and service in the field.

This is not a pivot away from what Alfa is. It is the next move in a line of moves: the early Markem distribution partnership, the in-house engineering capability that produced the Electronic Solutions portfolio alongside the Markem-Imaje relationship, and the integration of Aura's software capability into Alfa's engineering arm. Each of those moves was a step toward software as the center of the offer. This one names what the previous ones implied.

What the Books Say

The financial picture says it more clearly than any narrative can.

EGP **579**M TOTAL REVENUE FY2025 — the whole picture	EGP **428**M RECURRING REVENUE The Markem-Imaje annuity — inks, solvents, service	**74**% RECURRING SHARE The base this strategy protects
EGP **75**M ELECTRONIC SOLUTIONS — **FLAT SINCE 2023** — a product-feature gap	**2.16**% R&D INVESTMENT Today — funded like a distributor	**6–10**% R&D BENCHMARK Where software-led B2B invests

Of Alfa's EGP 579M FY2025 revenue, **EGP 428M — 74% — is recurring**, almost entirely from the Markem-Imaje annuity: inks, solvents, ribbons, spare parts, and the service relationships that keep them flowing. The Electronic Solutions division produced EGP 75M last year — and the year before that, and the year before that. The ES revenue line has been **flat since 2023**.

The recurring base looks healthy because Markem-Imaje as a global partner is healthy. The Electronic Solutions plateau looks like a sales problem because the sales force has not changed. Neither read is accurate.

The recurring base is exposed to Markem-Imaje's own AI evolution: the 24/7 AI-powered virtual assistant they launched in 2024, the new AI inspection software they introduced at Empack 2026, the IoT-connected services portfolio they have been building. As their MEA distributor, Alfa either evolves alongside them — capable of selling, installing, and servicing intelligent products — or becomes, over time, a courier of consumables to customers who are being courted directly. *The annuity is protected only as long as the partnership is.*

The Electronic Solutions plateau is not a sales problem. It is a product-feature problem. Q-MATIC has been shipping AI voice agents in production. Wavetec released a new AI forecasting module last month. Tensator has integrated facial recognition for passenger flow at airports. Customers walking into a tender in 2026 see vendor decks featuring AI capability across the queue management category, and Alfa's WAM platform — directionally right when it was built, faithfully maintained since — does not carry that capability. *The plateau is the market's quiet verdict.*

Underneath both of these realities is a third one. **Alfa invests 2.16% of revenue in R&D.** A distributor is funded that way; a software-led B2B company invests several times more. Moving from roughly 2% toward a 6–10% range over the next 24 months is the deliberately achievable first step — well short of where pure-software peers sit, but a credible, fundable correction of the gap between Alfa's R&D spend and Alfa's actual strategic identity.

The Bet

Over the next 24 months, Alfa will:

- **Protect the EGP 428M Markem-Imaje annuity** by becoming an AI-capable distributor — equipping the field service organization with AI-augmented tools, integrating Markem-Imaje's AI inspection and connected services into the local offer, and developing the customer-facing intelligence layer that turns consumable supply from a transactional relationship into a managed one.
- **Reinvent the Electronic Solutions line as a portfolio of intelligent software products on hardware substrates.** The queueing platform, the kiosk line, the hospitality and parking products, and the next generation of the smart utility line are reconceived from the inside: software as the center of the offer, hardware as the delivery substrate. The 2027 GS1 Digital Link inflection in the coding market creates a parallel upgrade wave Alfa must be positioned to lead among its existing customers, not lose to them.
- **Step R&D investment up from 2.16% toward a 6–10% range over the strategy window,** funded by the recurring base the investment is protecting. This is a deliberately conservative first move — the benchmarks for software-led B2B run higher — chosen to be achievable and fundable from a standing start, not aspirational. The investment is not speculative: it is the cost of staying in the business Alfa is actually in.

How

Execution runs through the engineering capability Alfa already has — Aura, operationally embedded as Alfa's software arm. A small AI Office is established to govern method, measurement, and data accessibility. Four to six AI Pods run at any given time, each owning a specific outcome with a measurable Phase 1 result by year-end 2026. A network of AI Champions in the divisions provides the connective tissue between pod work and day-to-day operations.

The first wave includes product and engineering, deliberately. Internal efficiency wins are useful, but the strategic asset is the AI Products line — and product cannot lead from a position of catching up. Phase 1 measures product outcomes and operational outcomes from day one.

The pacing is fast lanes inside a slow organization. Small, visible pods generate the wins that pull the rest of the company forward, rather than asking the company to adopt at the same pace as the pods.

Why Now

Three forces converge.

The 2027 GS1 Digital Link sunrise gives Markem-Imaje's customers an immediate reason to upgrade their coding and inspection systems — and gives Alfa's competitors an immediate reason to call those customers first.

The competitive AI gap in queueing is no longer emerging. The flat-since-2023 number tells the same story the competitive scan does, in a quieter voice.

And the partnership window with Markem-Imaje is the partnership window with Markem-Imaje. They are building the next generation of products on a known timeline. Alfa either matches that timeline or watches it pass.

✦

This document does not argue for AI adoption. The chairman is the company's first AI advocate, and the strategic intent is settled. The document describes the execution plan: what Alfa will build over the next 24 months, how the work will be organized, who will lead it, what will be measured, what will be discovered through the work, and how the strategy will be calibrated as evidence comes in.

> *"The rest of the document follows from this opening.*

SECTION 02

Vision — Alfa in 2028

By the end of 2028, Alfa operates as an AI-enabled software-on-hardware company. The brand presentation is still a recognizable Alfa — the same divisions, the same customers, the same Markem-Imaje partnership — but the substance of what is built, sold, and delivered has shifted. The center of every product is software intelligence; the hardware is the substrate that gives that intelligence physical presence in customer environments.

What follows is what that looks like in practice, told through the kinds of moments that would not be possible in 2025.

The Markem-Imaje base, intelligent

> A bakery in Alexandria runs three Markem-Imaje coders on its bread line. The plant manager opens an app on her phone in the morning. She sees that Coder #2 is operating at 92% of its expected character quality, down from 99% last week — Alfa's predictive print-quality model has detected ribbon-tension drift and scheduled a service visit for Thursday. She sees that her consumable usage is tracking 4% above forecast; the system has already adjusted next month's standing ink order. She does not think about Alfa at all. The system simply ran.

> In Riyadh, a multinational FMCG plant manager receives a message from her Alfa account team: a competitor coder vendor has filed a regulatory submission for AI-based code verification that may, in 18 months, give them an advantage on European supply contracts. Alfa is shipping the equivalent capability in Q1 2029. She forwards the message to her procurement lead.

In Cairo, an Alfa field service technician arrives at a packaging plant for a scheduled maintenance call. His tablet has already pulled the unit's last 90 days of operational telemetry, identified two parts likely to fail within the next 60 days, and pre-staged them in his truck. The visit takes 38 minutes instead of two hours. The customer's line never shuts down.

The AI Products line (formerly Electronic Solutions)

In a National Bank branch in Tanta, a customer walks in. The queueing system detects he is carrying a sealed envelope marked with a corporate logo and routes him to the SME relationship desk instead of the general teller queue — saving him 28 minutes on average. The system also notices that the SME desk is becoming a bottleneck this morning and quietly notifies the branch manager's screen, suggesting a temporary reallocation of two staff. None of this involves new hardware; the entire intelligence layer is software running on the existing displays.

A 4-star hotel in Sharm El Sheikh receives a guest at 11:43 PM. The Alfa hospitality platform pulls the guest's reservation history, recognizes a pattern of late check-ins, a previous request for a quiet room on a high floor, and an anniversary on the following day. The front desk associate sees this on her screen as the guest approaches. A complimentary upgrade and a small in-room gesture are already pre-cleared by the property manager's standing rules. The guest experiences a moment of being known.

A municipal smart-poles deployment in 6th of October City — ATUM v2 — adjusts its lighting output based on real-time foot-traffic patterns. It detects an unusual crowd gathering at 9:14 PM and pings the relevant municipal officer with a flagged image and a recommended response. Power consumption across the deployment is 31% below the original deployment baseline.

> A self-service kiosk in a public hospital in Mansoura takes vitals from a patient, asks a short conversational triage in Arabic, and produces a draft note for the nurse. The nurse reviews, edits, and accepts the note into the patient's record. The kiosk has reduced the triage queue by 40% during peak hours.

Operational intelligence — how the company itself runs

The engineering organization — a small core plus a flexible Aura cohort — runs four AI Pods at any given time. Each pod owns a specific outcome: one is shipping the next iteration of the predictive consumables model; one is integrating vision-based quality verification into a European OEM customer pilot; one is reducing the average proposal generation time across the sales floor; one is building the second generation of the municipal IoT platform. Pods rotate every 8 to 12 weeks. Champions in each department escalate problems worth a pod's attention; not all of them get one.

> A sales rep at the Tanta branch needs to respond to a tender from a regional retailer. She types a paragraph describing the customer's situation into a tool. Sixteen minutes later, she has a draft proposal with the right product configuration, defensible pricing, references to comparable customer deployments, and a properly formatted PDF in Arabic. She edits for tone and sends. Her predecessor in 2025 would have spent two days on the same task.

> The finance team closes the monthly books in three days instead of nine. The HR team can answer "which sales reps in our team have closed the most kiosk deals in healthcare verticals this year" in under a minute, without opening Excel. The customer service inbox is triaged automatically; only the small fraction of messages that actually require human judgment reaches a human.

The internal experience

Working at Alfa in 2028 feels different. New engineers join Alfa and assume that AI is part of the substrate of their work — they have never known a different way to build. Sales people who, in 2025, were trapped reselling queueing systems they understood, now sell software products that explain themselves through interactive demos in the customer's office. Field service technicians, who used to drive long routes hoping their truck had the right part, now show up with the right parts already loaded and a clear plan.

The chairman still walks the office in Cairo. The conversations he overhears are different. The engineering DNA he built into the company in 1987 is still here — but it has shifted forms. The company is still Alfa. It is also something that was not possible to be, before.

SECTION 03

The Strategic Choice

Three Workstreams, Sequenced Deliberately

The bet described in the Opening Hypothesis breaks into three workstreams. They are not three independent initiatives — they are three views of the same software-on-hardware transformation, sequenced so each feeds the next.

1 **Workstream 1: Annuity Intelligence.** Protecting the EGP 428M Markem-Imaje recurring base by becoming an AI-capable distributor. The work: embed AI-augmented tooling into the 75-person field service organization, integrate Markem-Imaje's AI inspection software and connected services into the local offer, and build the customer-facing intelligence layer that turns consumable supply from transactional into managed. The strategic logic: the annuity is exposed; defending it earns the credibility and the capability for the larger bet.

2 **Workstream 2: AI Products.** Reinventing the Electronic Solutions line from the inside. The queueing platform, the kiosk line, the hospitality and parking products, and the next generation of the smart utility line are reconceived with software as the center of the offer and hardware as the substrate. The strategic logic: the ES plateau is the market's verdict on the current product line; the plateau breaks when the product line catches up to where customer expectations have already moved.

3 **Workstream 3: Operational Intelligence.** Internal AI adoption across sales, finance, HR, customer service, and operations. The strategic logic: this is the rehearsal ground for the AI capability we are shipping in our own products. A company that cannot use AI internally cannot credibly sell intelligent software externally.

These three workstreams run in parallel from day one. The critical sequencing decision — and the one most likely to be mishandled in execution — is that **Workstream 2 (AI Products) cannot wait for Workstream 3 (Internal AI) to mature.** The familiar industry pattern of "we'll get good at AI internally first, then ship it externally" produces companies with productive sales teams shipping yesterday's products. Product and engineering lead the first wave, not the last.

The Two-Division Frame

Electronic Solutions and Marking & Packing have materially different customer profiles, sales motions, AI surfaces, and competitive threats. The strategy is **one unified framework with two execution tracks** — shared governance, shared learning architecture, shared methodology; division-specific tactics.

For the Markem-Imaje side, the AI surface is largely defined by the partner: AI inspection, IoT-connected services, predictive consumables management, the 2027 GS1 Digital Link transition. Alfa's role is to integrate, localize, and service these capabilities for the MEA customer base, with Arabic-first interaction as a structural differentiator the global partner cannot easily provide.

For Electronic Solutions, the AI surface is owned end-to-end. The queueing platform's predictive routing, the kiosk's conversational interface, the hospitality platform's guest recognition, the intelligent municipal layer of the smart utility line — these are Alfa product decisions, not partner integrations. The reinvention is structural, and the R&D investment step-up funds it.

Pacing — Fast Lanes Inside a Slow Organization

The pacing principle is deliberate and contrarian. The instinct in a hierarchical industrial company is to roll change out uniformly: announce the program, mandate participation, measure compliance. We are not doing this.

The strategy concentrates early momentum in small, visible **AI Pods** — four to six at a time, six to twelve weeks per pod, cross-functional, each with a specific outcome to ship. Pods run with executive air cover but minimal procedural friction. The wins they produce — a sales team closing more tenders because proposals are drafted by AI, a field service visit that takes 38 minutes instead of two hours, a queueing demo where the system answers customer questions in Arabic — generate organic pull from the rest of the company.

The alternative — mandating company-wide adoption from day one — is what produces the typical enterprise AI failure pattern: high compliance, low absorption, no shipped product. The decision is to let the rest of the company catch up at its own pace, while a fast lane delivers the strategic capability.

This is the synthesis of the top-down/bottom-up question. Top-down provides the permission, the budget, and the chairman's visible air cover. Bottom-up provides the velocity, the evidence, and the wins that make the air cover worth giving. Neither alone produces the outcome.

Decision Discipline — Do, Measure, Validate, Calibrate

Each pod operates on a Do-Measure-Validate-Calibrate loop. The bar to **do** is low: try things, ship working drafts, accept early imperfection. The bar to **validate** is rigorous: did this actually move the metric we said it would move? Critical challenge happens at validate, not at do — which is the inverse of how this kind of work is typically run in a careful organization, where critique kills initiatives before they ship.

The implication for governance: the AI Council does not gate experimentation. It gates progression — from validated experiments into production, and from production into broader rollout. A pod that has run for eight weeks and produced something the validate criteria approve becomes a candidate for scaling. A pod that has run for eight weeks and produced something the validate criteria reject is a learning, not a failure — and the next pod incorporates the learning.

This discipline matters because it inverts the cultural default. In a hierarchical organization, the instinct is to plan thoroughly before doing. In an AI-enabled software organization, the cost of trying is low and the cost of not trying is high. The methodology gives the chairman a way to see that experimentation is bounded and accountable, not chaotic — and gives the operating teams permission to try things without first writing a six-month plan.

This is also the synthesis of the course-based/problem-based learning question. The chairman's underlying goal — a company-wide AI ethos — is honored. The method that delivers it is real work organized as a learning system, with formal training filling specific gaps that surface in pod work, rather than a pre-packaged curriculum delivered to people who do not yet know what to do with it.

An Illustrative Case: What the Strategy Unlocks

One product in the existing portfolio illustrates the pattern this strategy is built to address — not because it is the centerpiece of the bet, but because it shows, in miniature, why a directionally sound product can fail for reasons that have nothing to do with the product itself. ATUM is used here as one example among several; the queueing platform, the kiosk line, and the hospitality products each tell a version of the same story.

ATUM was built in 2016, when the local market was still selling pure hardware. The product was directionally right: smart utility management, IoT-led, software-centered, ahead of the curve. It is also, in 2026, a zero-revenue line — no units deployed.

The instructive reading is this: ATUM in 2016 was a correct product intuition that the rest of the company was not yet structured to execute on. The sales force was a hardware-selling sales force. The company had not named software as its identity. Customers approached for it were being asked to buy something that did not match what they understood Alfa to be. The product was, in effect, early on its own company.

The lesson generalizes beyond any single product. The strategy this document describes is what makes a product like this viable: software named as Alfa's identity, sales and product reorganized around intelligent products, the engineering capability to ship and iterate, and Arabic-first AI as a structural advantage in regional deployments. Where such a product belongs in the eventual roadmap — and how much weight it should carry relative to the queueing, kiosk, and hospitality lines — is a decision for the divisional heads and the AI Council to make on the evidence, not a determination this strategy presupposes. It is named here as illustration, not as a Phase 1 commitment.

What We Are Not Doing

The cuts are part of the bet. The strategy does not include:

- **A six-month formal AI training course as the primary learning vehicle.** Capability is built through doing, not through curriculum. The learning architecture is embedded in pod work, with formal training filling specific gaps as they surface.

- **Company-wide AI adoption from day one.** Adoption follows visible wins, not mandates. Departments that pull AI in earn it; departments that wait can wait.
- **A wholesale data infrastructure rebuild as a prerequisite.** The current data accessibility constraint is real and elevated to a Phase 1 priority — but the pods do not stall while infrastructure is upgraded. They work with the data they can reach, and the data work runs in parallel.
- **Building AI capability from scratch where buying suffices.** Off-the-shelf foundation models, vendor APIs, and partner offerings are used wherever they fit. Custom engineering is reserved for the layers that differentiate Alfa's products.
- **Chasing every AI tool that emerges.** The strategy is product-and-outcome-led, not tool-led. New tools are evaluated against the workstream they would serve, not against their general impressiveness.
- **Treating internal efficiency wins as the success story.** The success story is the AI Products line shipping. Internal wins are how the company learns to ship.

Closing Frame

The choice described in this section is structural, not tactical. It is a choice about what kind of company Alfa is becoming, what it is willing to stop doing in order to become that, and what discipline it will hold itself to during the transition.

The next section — Execution Architecture — describes how the choice is implemented: the governance bodies, the pod structure, the champions network, the data and tooling layer, the cross-divisional mechanisms. It follows from this choice, not the other way around.

SECTION 04

Execution Architecture

Four Layers

The execution structure has four layers, each with a distinct purpose and a clearly bounded authority:

Sponsorship	Provides air cover, removes obstacles, names the priority. Holds the political weight.
Governance	The AI Council. Decides what scales, what stops, what merges into business-as-usual. Holds the strategic weight.
Execution	AI Pods (the doers) and the Champions Network (the connective tissue). Holds the delivery weight.
Enablement	The AI Office. Provides tooling, methodology, data accessibility, training. Holds the capability weight.

The layers are designed so that each has the resources and authority to do its job without depending on the others to make individual decisions. The Sponsorship layer does not approve pod work. The AI Council does not run pods. The pods do not negotiate budgets. The AI Office does not own strategic direction. This separation is deliberate — it is what makes the system run faster than a hierarchical command structure can.

The Governance Layer: The AI Council

The AI Council is the strategic body that governs the transformation. It is chaired by the Head of Software Development and includes the heads of both divisions, the Head of HR (representing the workforce capability dimension), the Head of Innovation, and one or two rotating cross-functional members on six-month rotations.

The deliberate composition signal: **the Council is not an engineering committee.** A clear majority of seats are non-engineering. The chair sits in software development because the methodology and infrastructure work originates there, but the strategic authority is collective. This is the explicit counter to the perception risk that the initiative is "an engineering thing." If the Council ever looks like a software department meeting, the structure has failed.

The Council meets every two weeks for the first six months of execution, then monthly once the operating rhythm has settled. Its authorities, in order of weight:

1. Approves which pods receive resources for the next quarter.
2. Decides which validated pods scale into production or merge into business-as-usual.
3. Approves the R&D budget envelope and its allocation across the three workstreams.
4. Resolves cross-divisional dependencies and conflicts.
5. Acts as the formal interface to the Sponsorship layer — keeping the chairman informed without requiring the chairman's involvement in routine decisions.

The Council does not run pods, write code, or approve technical implementation choices. Pods come to the Council with validate-criteria results. The Council says scale, stop, or extend.

The Enablement Layer: The AI Office

The AI Office is small by design — two to three people initially, growing to five or six by year-end of Phase 1. It sits inside the engineering organization for reporting and infrastructure purposes but serves the entire company. Its scope:

- **Tooling and infrastructure.** Standing up the AI development environment, securing vendor and model access, managing API keys and credentials, providing the common substrate that pods build on.
- **Methodology custodianship.** Owning the Do-Measure-Validate-Calibrate playbook. Helping pods scope their validate criteria. Capturing learnings across pods so the next pod inherits the wisdom of the last.
- **Data accessibility.** Elevated to a co-equal priority for Phase 1. The financial data exercise that built the strategy's baseline demonstrated that existing data infrastructure cannot, today, answer questions a software-led B2B company needs to answer routinely. The AI Office owns the multi-quarter program to fix this — not by replacing systems wholesale, but by building the access layer pods need to ship.
- **Training delivery on demand.** Curating and delivering targeted training when pod work surfaces specific capability gaps. Not running a curriculum.

The AI Office does not own strategy. It does not run pods. It is a services organization for the rest of the structure.

The Execution Layer: AI Pods

A pod is a cross-functional team of four to seven people, time-boxed to six to twelve weeks, dedicated to a specific outcome with a measurable validate criterion. The anatomy:

- **A pod lead.** Owns the outcome. Reports to the AI Council on progress and validate-criteria results. Typically a senior contributor from the function whose outcome the pod is improving — not always an engineer.

- **One or two engineers from Aura.** Provide the technical capability.
- **Two or three operational contributors from the relevant departments.** Provide the domain context, sit closest to the customer or process being improved, and become embedded champions when the pod ends.
- **A periodic check-in from the AI Office.** On methodology, not on decisions.

The pod's first deliverable is a written **validate criterion** — a single, measurable claim the pod is testing. Examples: *"Field service visit time on coder maintenance calls drops by 30% with AI-augmented diagnostic tooling." "Sales proposal generation time drops from two days to under one hour with AI drafting." "Predictive consumables forecasting reduces standing-order variance by 25%."* The criterion is approved by the AI Council before the pod starts work.

The pod's lifecycle has three checkpoints:

- **Week 2.** Initial demo of working draft. The Council confirms the pod is on the trajectory it described.
- **Week 6.** Mid-pod validation read. The pod presents what is and is not working. The Council may extend, narrow, or stop.
- **Week 8–12.** Final validate read. The Council decides: scale, merge into business-as-usual, stop with a learning, or extend for one additional iteration.

Four to six pods run simultaneously. They are spread across the three workstreams so that the company is making progress on Annuity Intelligence, AI Products, and Operational Intelligence in parallel — not sequentially. Each pod is publicly visible across the company. Their wins are shared. Their stops are explained.

The Execution Layer: Champions Network

The Champions Network is the connective tissue between pod work and day-to-day operations. One to two champions per department, named at Phase 1 launch, with a fraction of their time formally allocated to AI champion duties.

Champions are not pod members by default — they are the people who:

- Surface AI-relevant problems from inside their department to the AI Council.
- Receive pod outputs and integrate them into the department's workflow when a pod completes.

- Train colleagues in their department on tools and methods that have been validated and rolled out.
- Provide ground-truth feedback to the AI Office about what is and is not working.

The Champions Network is the mechanism by which **pull replaces push**. Departments that have an active champion get AI capability first; departments that do not, do not. Champions are selected from inside the department, not assigned from outside it — which means an engaged department can effectively volunteer itself for the next wave of pod work, and a disengaged department defers itself.

Two implications worth flagging:

- **Champions are explicitly distributed across functions, not concentrated in engineering.** Sales, field service, finance, HR, customer service, and both divisions' operations each have a champion or two. Engineering is included but is one of many.
- **Champions earn their role through interest and effort.** The first wave is selected by the AI Council; subsequent waves grow organically from people who have engaged with pod work.

The Data and Tooling Substrate

The execution structure sits on a substrate of tooling, model access, and data accessibility. The financial data exercise that built the strategy's baseline made this visible: the existing infrastructure cannot, today, answer basic questions a software-led B2B company needs to answer routinely. The fundamentals took weeks to extract.

The substrate work is therefore not optional and not deferred. It runs as a dedicated workstream owned by the AI Office, with a multi-quarter horizon. The work is bounded by a single principle: **the substrate serves the pods, not the other way around.** Pods do not stall waiting for clean data. They work with what they can reach and flag what they cannot. The substrate team prioritizes accessibility work based on what the pods are actually blocked by — a feedback loop that prevents the common failure pattern of multi-year data warehouse rebuilds with no business value attached.

In practical terms, the substrate work in Phase 1 includes:

- Standing up the central AI development environment (model access, vendor APIs, key management).
- Building a common data-access layer over the Dynamics AX financial and operational data, prioritized by what pods need.

- Establishing the customer and service data flow that the Annuity Intelligence workstream depends on.
- Creating the shared knowledge base into which pod learnings accumulate.

Cross-Divisional Mechanisms

Electronic Solutions and Marking & Packing have different customer profiles, different sales motions, different competitive contexts, and — crucially — different employee bases that interact with each other less than is healthy for shared learning. The execution structure includes explicit mechanisms to keep the framework unified without either division dominating:

- **AI Council composition** ensures both divisions have a voice at the strategic table.
- **Pod composition** is required to span divisions where the outcome warrants it. A pod working on AI-augmented field service touches both divisions' field operations and is staffed that way.
- **Champions Network** is balanced across both divisions, with quarterly cross-divisional gatherings to share learnings.
- **The shared AI Office** serves both divisions, with no parallel division-specific capability allowed to grow.

The intent: the divisions retain their distinct execution cultures, but the AI capability is a shared company asset, not a per-division initiative that could fork.

Cadences

The operating rhythm has a specific shape:

- **Weekly.** Each pod runs its own internal cadence (typically a working session and a status update). The AI Office holds an open clinic that any pod can attend for methodology or substrate questions.
- **Every two weeks (first six months) or monthly (after).** AI Council meeting. Pod check-ins, scale/stop/extend decisions, cross-pod learnings, strategic adjustments.

- **Monthly.** Sponsorship layer briefing. A short, visible update to the chairman with progress against the strategy's spine — what is being learned, what is shipping, what has stopped, what has surfaced.
- **Quarterly.** Strategy review. The AI Council and Sponsorship together review the trajectory, calibrate the strategy against accumulated evidence, and reset the pod portfolio for the next quarter.
- **Half-yearly.** Phase review. Did Phase 1 (or 2, or 3) achieve what it was meant to? Does the strategy still hold? What is the next phase?

The cadences are designed so that fast pod work happens inside a stable governance rhythm. Pods iterate weekly; strategy adjusts quarterly; the spine moves half-yearly. The chairman sees a monthly heartbeat without being asked to make weekly decisions.

Closing

The architecture described here is the machine that delivers the choice described in the previous section. It is small enough to be fast, structured enough to be accountable, and distributed enough to avoid being captured by any one function.

The next section — Evidence and Reasoning — describes why this shape is the right one, drawing on what comparable enterprise AI adoption efforts have learned the hard way, and on the financial and competitive realities documented earlier in this document.

SECTION 05

Evidence & Reasoning

This section makes the case for the specific shape of the strategy described in the preceding sections — the three workstreams, the product-first sequencing, the AI Pods with validate criteria, the Champions Network as a pull mechanism, the elevation of data accessibility to a co-equal Phase 1 priority. Each of these choices is supported by what comparable organizations have learned in their own AI adoption efforts, what the financial data on the table demands, and what is established about how new capability actually moves through an organization.

The argument is not *"AI is important."* That argument is settled. The argument is *"this shape is the right shape, for this company, at this moment."*

What Comparable Organizations Have Learned

A recognizable pattern has emerged across enterprise AI adoption between 2023 and 2026. The pattern has both a positive and a negative form.

The positive form appears in organizations that have shipped real AI capability — including some of the most-studied enterprise rollouts. The common elements:

- Visible executive sponsorship without executive micromanagement. The sponsor names AI as a priority, removes obstacles, and shows up at moments that matter — but does not approve individual experiments.
- Small, time-boxed teams owning specific outcomes, rather than centralized AI teams owning everything. The decentralization is structural, not incidental.
- Measurable claims at the start of every initiative — what would success look like, in numbers, with a deadline. Initiatives that cannot articulate this do not get resourced.

- Cross-functional composition of working teams — engineering paired with domain experts, never engineering working alone for non-engineering users.
- Early publication of both wins and stops. Stops are normalized; they accumulate as learning rather than embarrass.

The negative form is observed in organizations that have run high-budget AI programs for two or three years and not shipped meaningful AI capability into their products or operations. The common failure modes:

- The initiative is captured by a single function — usually IT or engineering — and the rest of the company treats it as someone else's project.
- A centralized AI team is established and accumulates ownership of everything, becoming a bottleneck rather than an enabler.
- The methodology becomes talk-heavy: workshops, training programs, communities of practice, and conference talks accumulate without anything shipping to a customer.
- Internal efficiency wins are celebrated and the AI Products line never materializes. The company gets good at writing email faster and never ships an AI-enabled product.
- The data infrastructure work becomes a multi-year prerequisite that prevents any actual AI work from starting.
- One high-profile customer-facing launch is rushed without the validate discipline, performs poorly, gets walked back publicly, and becomes a cautionary tale that slows internal adoption for the following year.

The strategy in this document is shaped by both forms. The positive pattern shows up as the AI Pod model with validate criteria, the cross-functional Council, the Champions Network as connective tissue, and the chairman's visible-but-not-micromanaging sponsorship role. The negative pattern shows up as everything we have explicitly committed *not* to do — the centralized AI team, the months-long curriculum as the primary learning vehicle, the data rebuild as prerequisite, the internal-first product-later sequencing.

The reasoning is not *"Company X did this"* or *"Company Y did that."* It is that the field has now had three years of public attempts, the patterns have stabilized enough to learn from, and the strategy is calibrated against those patterns.

The Financial Logic of the Bet

The financial argument rests on three claims, each grounded in numbers on the table.

Claim 1: The current R&D investment is structurally incompatible with the company's strategic identity.

Alfa invests 2.16% of revenue in R&D — EGP 12.5M against a 579M revenue base. This is the R&D profile of a distributor. Distributor norms run roughly 2–3% of revenue, and the spend is largely operational — sales tools, vendor-supplied product configuration, partner-driven training.

Software-led B2B companies — which is what the strategy commits Alfa to becoming — invest far more. Independent benchmarking surveys of more than a thousand private B2B software companies place median R&D investment in the high teens to low twenties as a percentage of revenue (see References). Alfa is not proposing to reach that level inside this strategy window, and it would not be credible to claim otherwise from a standing start of 2.16%.

The target is therefore deliberately set below the benchmark: a step from roughly 2% toward a 6–10% range over 24 months. This is large enough to fund a genuine product-development capability and small enough to be fundable from the existing recurring base without straining it. The point is not to match pure-software peers in the first phase — it is to stop funding a software-led identity at a distributor's rate. The gap this closes is the gap between what Alfa says it is becoming and what its investment actually supports.

Claim 2: The cost of evolving the partnership is structurally lower than the cost of losing the annuity.

The EGP 504M Markem-Imaje contribution to revenue is 87% of the top line. Of that, EGP 428M is recurring. Recurring revenue at this scale is the cheapest revenue to protect and the most expensive revenue to lose.

Markem-Imaje's own evolution toward AI-mediated services — the 2024 24/7 AI assistant, the 2026 AI inspection software, the IoT-connected services portfolio — creates a partnership-readiness obligation for Alfa as their MEA distributor. The cost of evolving the local field service organization, integrating their AI tooling, and developing the customer-facing intelligence layer is in the low single-digit millions of EGP per year over the strategy window. The cost of being replaced as their MEA partner, or of being marginalized while a more capable distributor is courted, is the annuity itself.

The math is asymmetric. A small investment protects a large recurring stream. Even modest probability of partnership-stress in the absence of the investment justifies the spend.

Claim 3: The Electronic Solutions plateau is unlockable, and the unlock pays for the work that unlocks it.

The ES line has been flat at roughly EGP 75M since 2023. Queueing systems account for 64% of that. The competitive scan documents the reason: customers in 2026 are seeing AI-enabled queueing systems from Q-MATIC, Wavetec, and others. Alfa's WAM platform does not carry that capability. The plateau is the product-feature gap.

A successful re-base of the queueing platform with AI features — predictive routing, voice or chat agents, natural-language analytics, Arabic-first interaction — moves the addressable opportunity per tender materially. A return to growth that recovers even the 2018–2022 trajectory adds tens of millions of EGP to the ES line within the strategy window. The R&D investment funds itself within two to three years, and continues compounding after.

> The three claims together describe a **self-funding transformation:** the recurring base pays for the partnership evolution, the partnership evolution preserves the recurring base, the ES re-base unlocks growth, and the growth pays for further reinvestment. The strategy is not a cost program. It is a reallocation from underspending on identity to investing at the rate identity demands.

What the Competitive Scan Says, in One Paragraph

The full scan is in the appendix. The summary that matters here: in both of Alfa's core markets, AI capability has moved from "emerging" to *operational and at scale*. Q-MATIC ships AI voice agents, Wavetec ships predictive forecasting and natural-language analytics, Tensator has integrated facial recognition for passenger flow. In industrial coding, Domino has shipped more than a thousand AI-vision inspection units since 2020 and launched a printhead purpose-built for the 2027 GS1 Digital Link transition. Markem-Imaje is actively closing the gap — AI inspection software, 24/7 AI service tools, IoT-connected printers. The window for entering the AI products line as a differentiated player closes in 2027. Entering as a follower happens automatically; entering as a leader requires the action this strategy commits to.

How New Capability Actually Moves Through Organizations

The Champions Network and the pull-based adoption model are not arbitrary design choices. They follow from a well-established pattern in how innovation spreads inside organizations: not as a uniform wave, but in a sequence of distinct adopter groups.

A small fraction of people in any organization are oriented toward trying new tools immediately — they will engage regardless of the program's design. A larger early-adopter group will engage if the work is visible, credible, and supported by people they respect. The early majority will engage when the early adopters' wins become socially evident — when colleagues they trust have visibly benefited. The late majority will engage when *not* engaging starts to feel like falling behind. And a small group will hold out indefinitely, often productively as a counterweight to over-enthusiasm.

Programs that try to move everyone at the same pace produce uniform low engagement. Programs that concentrate effort on the early waves and let the later waves pull themselves in produce uneven but real adoption. The Champions Network is the implementation of this pattern: it identifies and resources the early adopters, makes their work visible, and lets the social mechanism do the work that no mandate can do.

The same pattern shows up in what is known about how adults build new capability. New skills are built mostly through doing real work — roughly seventy percent of the learning. Another twenty percent comes from interaction with skilled peers. Only the last ten percent comes from formal instruction. The pod model, the Champions Network, and the on-demand training delivery from the AI Office together honor this distribution. A six-month formal curriculum delivered to people who are not yet doing the work inverts it.

What Skilled AI Use Looks Like

The strategy's commitment to embedded learning has a counterpart in what "skilled AI use" actually means at the individual level. Across enterprise practice, four capabilities are emerging as the consistent markers of someone who uses AI well:

Delegation	Knowing what to hand to an AI and what to keep. The judgment is task-by-task, not technology-by-technology.
Description	Being able to articulate what is needed, with the right context, at the right level of specificity. The bottleneck is rarely the AI; it is the human's ability to describe the work.
Discernment	Evaluating AI output critically before acting on it. The skill that prevents the AI-generated mistakes that erode trust.
Diligence	Maintaining responsibility for outcomes, including outcomes the AI participated in. AI does not absolve.

These four capabilities are not taught well in courses. They are built through repeated cycles of trying, failing visibly, and recalibrating with peers. The pod model is designed to produce these capabilities as a side effect of shipping outcomes, rather than as the primary objective of training programs.

Evidence the Strategy Will Generate

Phase 1 produces its own evidence base. The validate-criteria results from each pod, the patterns of which departments engage and which defer, the substrate work's actual cost and pace, the customer-facing pilots' reception in the two divisions' verticals — these accumulate into the data the strategy uses to calibrate itself. Phase 2 begins with this evidence in hand and adjusts accordingly.

The strategy does not require all of its bets to be right on day one. It requires the validate-and-calibrate discipline to be honored — which is why the methodology occupies as much space in the document as the structural choices do.

Closing

The reasoning behind the strategy is the same as the reasoning behind any considered bet: the trade-offs have been examined, the costs and benefits have been weighed against the data available, the patterns from comparable efforts have been incorporated, and the structure has been designed to learn and adjust as new evidence arrives.

The next section — People and Sponsorship — describes the human structure that holds this together: who plays which role, how the sponsorship transitions over time, and how the political weight of the transformation is distributed across the people who will own the outcomes.

SECTION 06

People & Sponsorship

The strategy is held together not by a single owner but by a distributed structure of stewardship. This section describes who holds which role, how the sponsorship transitions over the strategy window, and how the political and operational weight of the transformation is distributed deliberately rather than concentrated.

The discretion convention applies throughout: roles are described in structural and functional terms, by the seat they hold. Individual assignments are intentionally deferred. At this stage the strategy defines the seats and their authorities; who fills each seat is a decision for the chairman and the AI Council to make once the strategy is endorsed. Keeping names out of the document at this stage is deliberate — it keeps the focus on the structure and the tactics, and avoids prematurely anchoring the design to specific individuals.

The Sponsorship Architecture

Four layers, each with a distinct kind of authority and a distinct kind of accountability:

- **Sponsorship layer.** Holds the political weight of the transformation. Names the priority, removes obstacles, makes the work visible inside the company and to external partners. Operates at the rhythm of monthly briefings and quarterly reviews — not weekly decision-making.
- **Strategic governance layer.** The AI Council. Holds collective strategic authority over what scales, what stops, and what merges into business-as-usual. Composed deliberately across functions so that no single department's agenda dominates.

- **Operational leadership layer.** The Strategy Leader (the Council chair) and the workstream leads. Hold day-to-day responsibility for execution quality across the three workstreams.
- **Strategy continuity layer.** The Strategy Architect role, evolving from active drafting into advisory presence over the strategy window. Holds the institutional memory of why specific structural choices were made and the discipline to surface drift when it begins.

The layers are not a hierarchy in the traditional sense. The Sponsorship layer does not direct the Strategic Governance layer's decisions. The Strategy Continuity layer advises across the others without commanding any of them. The Operational Leadership layer reports to the Strategic Governance layer on results, not on plans. This structure is designed for durability: any single departure from any layer should not derail the transformation.

The Sponsorship Layer in Phase 1

For the opening 6–12 months of execution, the chairman holds the visible sponsorship role. In practice this means:

- Presence at the Phase 1 kickoff and quarterly milestones, with brief, direct remarks framing the strategic priority.
- Active naming of AI as the company's transition theme in leadership meetings, customer presentations, and partner conversations — particularly with Markem-Imaje, where the partnership-readiness signal lands at the executive level.
- Removing organizational obstacles when surfaced — budget approvals, headcount additions, cross-divisional conflicts that need senior intervention.
- Receiving a monthly heartbeat briefing from the AI Council: what is shipping, what is stopping, what is being learned, what needs attention.

What sponsorship explicitly does *not* include:

- Approving individual pod work or technical decisions. The Council holds that authority collectively.
- Attending every Council meeting. The cadence is a monthly briefing, not bi-weekly attendance.
- Direct intervention in pod execution. Visible support, yes; operational intervention, no.

This boundary is the political mechanism that lets the structure move at the speed the strategy requires. A sponsor who is visible but not interventionist is the rarest and most valuable form.

The Sponsorship Transition

The sponsorship layer is designed to transition from the chairman to the heads of the two divisions over the strategy window. **The transition is condition-driven, not date-driven.**

> **TO BE DEFINED WITH THE CHAIRMAN**
>
> The specific transition conditions are a decision for the chairman and the divisional heads. Working framing for that discussion: the transition completes when (a) Phase 1 has produced a baseline of validated results visible to the divisions, (b) the divisional heads have formally co-sponsored at least one pod each across the workstreams that touch their division, and (c) the divisional heads are present at quarterly reviews as principals, not observers. The chairman remains available as an escalation path but no longer holds the visible sponsorship role.

The reason the transition is conditional rather than scheduled: a date-driven handoff risks transferring sponsorship before the divisional heads have built the relationship with the work that gives the sponsorship credibility. A condition-driven handoff transfers sponsorship when it is genuinely owned, not when the calendar says it should be.

The Strategy Leadership Layer

The Strategy Leader chairs the AI Council and holds operational responsibility for the transformation's execution quality. This seat sits in the Head of Software Development role because the methodology, the engineering capability, and the infrastructure substrate all originate from that organization.

This placement carries an explicit risk identified earlier in the strategy: that the initiative could be perceived as "an engineering thing" by non-technical functions. The structural counter is the AI Council's composition, which gives non-engineering seats a clear majority. The political counter is the workstream lead structure described below, which distributes operational ownership across functions.

The Strategy Leader's specific authorities:

- Chairs the AI Council and shapes its agenda.
- Coordinates across the three workstream leads.
- Owns the relationship with the AI Office and the engineering capability supplied by Aura.
- Represents the transformation to the Sponsorship layer in the monthly heartbeat briefing.
- Surfaces strategic drift or under-investment to the Strategy Continuity layer.

The Governance Body — AI Council Composition by Role

The AI Council seats are composed deliberately to distribute strategic authority across functions:

- **Chair.** The Strategy Leader.
- **Head of Electronic Solutions division.** Holds the operational view of the AI Products workstream as it touches the ES portfolio. ⊘ Seat — individual not yet assigned.
- **Head of Marking & Packing division.** Holds the operational view of the Annuity Intelligence workstream as it touches the Markem-Imaje base. ⊘ Seat — individual not yet assigned.
- **Head of HR.** Holds the workforce capability dimension of the transformation — the human side of how new capability moves through the company. Engaged from day one through a review-and-shape loop with the strategy as it develops.
- **Head of Innovation.** Holds the external-facing surface of the transformation — partnerships, scouting, and the connection between Alfa's AI direction and the broader ecosystem of vendors, research partners, and customer segments. Engagement staged in after the strategic spine is established.
- **Rotating cross-functional voices.** One or two seats on six-month rotations, drawn from functions whose work is most relevant to the current pod portfolio (sales, field service, customer success, finance, depending on the moment).

The Head of HR seat warrants a specific note. The workforce capability dimension — how the people in the company actually build AI capability, how roles evolve, how the learning architecture lands in the HR systems and the company culture — is the layer most likely to be neglected by an engineering-led program and the layer most consequential for whether the transformation sticks. The Head of HR's day-one involvement is not a nominal seat; it is structurally load-bearing.

Workstream Stewardship

Each of the three workstreams has a designated lead who holds operational responsibility for the workstream's pods and outcomes. A fourth lead role covers the external-facing partnership and scouting work that connects the three workstreams to the broader ecosystem.

- **Annuity Intelligence lead.** Operational stewardship of the Markem-Imaje partner-readiness work — field service tooling, AI inspection integration, customer-facing intelligence. ⊘ Individual not yet assigned; natural fit is within the Marking & Packing division or a senior figure within field service operations.
- **AI Products lead.** Operational stewardship of the Electronic Solutions reinvention — queueing, kiosks, hospitality, parking, and the next generation of the smart utility line. ⊘ Individual not yet assigned; natural fit is within the Electronic Solutions division in close coordination with Aura engineering.
- **Operational Intelligence lead.** Stewardship of the internal AI adoption work across sales, finance, HR, customer service, and operations. ⊘ Individual not yet assigned; natural fit is the Head of HR seat, given the workforce capability overlap.
- **External Partnerships and Scouting lead.** Stewardship of the partnership, ecosystem, and external-surface work. ⊘ Individual not yet assigned; natural fit is the Head of Innovation seat.

The workstream leads do not chair the AI Council. They are accountable to it. The separation between Council governance and workstream operational leadership is intentional — it prevents any single workstream from capturing strategic direction.

The Strategy Architect's Transition

The Strategy Architect role — the originator of this document and the discipline behind its choices — is designed to transition from active drafting into an advisory presence over the strategy window:

- **Phase 1 (Q3 2026 — year-end 2026).** Active engagement: shaping the Council's agenda from outside, attending key milestones, working closely with the Strategy Leader on early validate-criteria results, supporting the chairman's sponsorship presence.
- **Phase 2 (early 2027 onward).** Advisory engagement: quarterly check-ins with the Strategy Leader, ad-hoc availability for strategic recalibration, surfacing drift when it appears.
- **Phase 3+ (late 2027 and beyond).** Continuity-of-memory role only. The operational ownership has fully transferred; the architect's role is to remember why specific choices were made and to surface them when context for those choices is needed.

> TO BE DEFINED WITH THE CHAIRMAN
>
> The specific timing of each transition and the conditions that mark each shift are a decision for the chairman, the Strategy Leader, and the Head of HR — given the transition's interaction with the workforce dimension.

The reason for explicit transition rather than indefinite continuation: a strategy that depends indefinitely on its original architect has not been institutionalized. The architect's job is to make themselves progressively unnecessary to the work continuing.

What This Distribution Achieves

The political weight of the transformation is distributed across distinct seats: the chairman holds visible sponsorship; the Strategy Leader holds operational leadership; the Head of HR holds the workforce dimension; the two divisional heads hold the eventual long-term sponsorship stewardship and the divisional execution surfaces; the Head of Innovation holds the external surface; the Strategy Architect holds the continuity-of-memory role in an advisor capacity.

> No single seat can capture the transformation. No single departure can derail it.

This distribution is itself the most important political choice the strategy makes. A program with a single owner — however capable — is exposed to that owner's availability, conviction, and tenure. A program distributed across the structure described here is owned by the company.

The next portion of this document — the Phase 1 Playbook appendix — describes the specific pods, the substrate work, and the first-90-days execution detail that brings this structure to life.

SECTION 07

Risks, Metrics, Calibration

The strategy is a hypothesis. This section describes what will be measured to know if the hypothesis is holding, what could cause it to fail, and how the strategy will adjust as evidence accumulates.

The framing matters: the strategy does not assume its own correctness. It assumes the methodology — Do-Measure-Validate-Calibrate — and trusts that the methodology will surface where the strategy is right and where it needs adjustment. The metrics, risks, and calibration mechanisms described here are the means by which that adjustment happens.

Metrics

The metrics fall into four groups: three mirroring the workstreams, plus a fourth covering cross-cutting strategic health.

Annuity Intelligence

The work is judged by whether the Markem-Imaje partnership becomes more resilient through Alfa's evolution.

Field service efficiency. Average time per coder maintenance visit.	Target: 30% reduction by year-end Phase 1
Consumables forecast accuracy. Standing-order variance for the top 50 customer accounts.	Target: 25% improvement in variance by year-end

Markem-Imaje partner signals. Qualitative indicators: invitation to early partner programs, references to Alfa in MEA-region presentations, technical integration depth. Tracked through the AI Council's quarterly partner-engagement review.

Customer retention proxies. Renewal rate and contract value evolution for the top 100 customers. Baseline established Q3 2026.

AI Products

The work is judged by whether the Electronic Solutions plateau breaks.

Tender win rate on AI-enabled bids. Tracked from the first tender in which the new AI-enabled queueing demo is presented.

Target: visible improvement vs. baseline by mid-Phase 2

Time from product idea to first shipped iteration. Measured against the 6–12 week pod timebox.

Target: 80%+ of pods ship within their timebox

ES revenue trajectory. Quarter-over-quarter ES revenue. Phase 1 succeeds if the line stops declining; Phase 2 succeeds if it grows.

Demo-to-discovery conversion. Percentage of customer demos that convert to pilot agreements. Baseline established Q3 2026.

Operational Intelligence

The work is judged by whether internal capability accumulates in a way that supports product evolution.

Time saved per role per week. Self-reported surveys plus calibrated task timings.

Target: 4+ hours/week for adopted roles by year-end

Adoption rate across departments. Percentage of department personnel actively using validated AI tools. Tracked via the AI Office.

Sales proposal generation time. From paragraph-of-context to formatted PDF.

Target: under 1 hour, down from baseline 1–2 days

Monthly close cycle. From 9 days to 5 days by year-end Phase 1.

Strategic Health

Cross-cutting indicators that signal whether the strategy itself is on track.

R&D as percentage of revenue, trajectory. From 2.16% baseline toward the 6–10% target. Reviewed quarterly.

Pod validate-criteria pass rate. Percentage of pods that deliver against their stated criteria. Pass rate above 70% suggests pods are well-scoped; below 50% suggests scoping is too ambitious or substrate constraints are too tight.

AI Council decision velocity. Average time from pod proposal to scale/stop/extend decision.

Target: under 30 days for routine decisions

Champions Network density. Number of active champions, distribution across functions and divisions, organic growth rate.

Workstream balance. Resource allocation across the three workstreams. **Watch indicator:** if Workstream 2 (AI Products) drops below 35% of pod allocation, the strategic risk identified at project start is materializing.

Risks

Eight risks are tracked actively. Each has a named mitigation and a leading indicator.

1

Internal efficiency wins crowd out product investment

The strategic risk identified at the project's start. The familiar enterprise pattern: efficiency wins are celebrated, product evolution stalls.

MITIGATION Product-side metrics tracked from day one; minimum two AI Products pods running at all times; quarterly review of workstream balance.

LEADING INDICATOR Workstream allocation drift toward Operational Intelligence.

2

"Engineering thing" perception captures the initiative

The cultural risk in placing the Strategy Leader inside software development.

MITIGATION AI Council composition with non-engineering majority; Champions Network distributed deliberately across functions; pod composition required to include domain experts from non-engineering departments.

LEADING INDICATOR Champion distribution skew toward engineering, Council attendance patterns becoming software-department-heavy.

3

Substrate work becomes a multi-year delay

The common failure pattern of treating data infrastructure as a prerequisite.

MITIGATION Substrate-serves-pods principle; pods proceed with what they can reach; substrate prioritized by pod blockers, not abstract architecture.

LEADING INDICATOR Percentage of pods reporting substrate blockers in their week-2 check-in.

4

Markem-Imaje partnership accelerates beyond Alfa's readiness

The partner may move faster than Alfa's evolution can match, leading to partnership stress or marginalization.

MITIGATION Annuity Intelligence workstream begins day one; named pod for partner AI integration; quarterly partnership review with explicit readiness assessment.

LEADING INDICATOR Partner-led conversations about other regional distributors; delayed Alfa involvement in new product rollouts.

5

2027 GS1 Digital Link transition catches Alfa underprepared

The specific industry inflection that creates exposure for both partner and direct customer relationships.

MITIGATION Dedicated tracking within the Annuity Intelligence workstream; coordinated with Markem-Imaje's own transition roadmap; customer-by-customer readiness assessment in 2026.

LEADING INDICATOR Customer questions about 2D code readiness that Alfa cannot confidently answer.

6

Sponsorship transition timing wrong

Either too early (no validated results to anchor divisional ownership) or too late (chairman bottleneck slows progression).

MITIGATION Condition-driven transition framework with explicit triggers; quarterly check on transition readiness.

LEADING INDICATOR Chairman becoming the deciding voice on operational matters that should sit in the Council.

7

Talent retention as Aura engineers see scaled opportunity

Engineers becoming AI-fluent in a market where AI talent is in high demand.

MITIGATION Visible scope and ownership of named workstreams; R&D budget growth that creates career room; explicit upgrade of titles and compensation as capability grows.

LEADING INDICATOR Voluntary attrition rate; exit interview themes.

8

Customer trust loss from poorly-shipped AI features

The rush-to-market pattern that has cost other enterprises publicly.

MITIGATION Validate criteria for every customer-facing shipment; controlled pilot rollouts; no big-bang launches.

LEADING INDICATOR Customer complaint patterns on AI-enabled features specifically.

Calibration

The strategy is itself a hypothesis with validate criteria. The calibration mechanism operates at three timescales.

Quarterly	The AI Council reviews accumulated pod validate-criteria results, adjusts the next quarter's pod portfolio, and refreshes the metrics dashboard with new data. Quarterly adjustments are tactical — which pods, what staffing, what substrate priorities.
Half-yearly	The Council and Sponsorship layer together review the strategic spine. The questions: Do the workstream priorities still hold? Are the financial claims trending as expected? Are competitive conditions evolving as the scan predicted? Half-yearly adjustments are structural — workstream weighting, governance composition, R&D allocation.
Annual	Full strategy review. Is the bet still the right bet? Is the identity claim — AI-enabled software-on-hardware company — still the right framing? Annual adjustments are foundational; they may revise the strategic identity itself if accumulated evidence warrants.

The strategy's correctness is not assumed at any timescale. The methodology is. The shape of the strategy will evolve; the discipline of validate-and-calibrate is what makes that evolution rigorous rather than reactive.

Closing

A strategy that cannot be measured cannot be steered. A strategy that does not track its own risks accumulates them silently until they surface as crises. A strategy that assumes its own correctness becomes brittle.

The metrics, risks, and calibration framework described here are not bureaucratic overhead. They are the mechanism by which the strategy stays honest with itself as the company executes against it.

APPENDIX

Phase 1 Playbook

Phase 1 is the first 90–180 days of execution, starting Q3 2026 and running through approximately year-end. Its purpose is to:

1. Launch the first wave of AI Pods across the three workstreams.
2. Stand up the substrate (AI Office, tooling environment, data accessibility) that subsequent pods depend on.
3. Produce a baseline of validated results that justify Phase 2 scaling.
4. Surface the first wave of strategic learnings that calibrate the broader strategy.

Phase 1 is deliberately bounded. It is not the strategy — it is the credible start of the strategy. The work that anchors months 6–24 of the bet is informed by Phase 1's evidence, not pre-determined by Phase 1's plan.

Proposed Phase 1 Pod Portfolio

Five pods proposed for the first wave, distributed across the three workstreams. The portfolio is sequenced for measurable validate-criteria results within Phase 1.

Pod 1A — Field Service AI-Augmented Diagnostics

OUTCOME	Reduce the average Markem-Imaje field service visit time through AI-augmented diagnostic tooling and pre-staged parts recommendation.
VALIDATE CRITERION	Average time per coder maintenance call drops by 30% in a pilot territory (Cairo or Alexandria) by week 10.
ANATOMY	– Pod lead: Senior figure from field service operations – 1 Aura engineer – 1 spare parts / logistics person – 1 Markem-Imaje technical liaison (if available)
TIMELINE	10 weeks.
SUBSTRATE DEPENDENCIES	Field service ticket history; equipment telemetry where available; Markem-Imaje technical bulletins.
WHY FIRST WAVE	Direct, measurable impact on a 75-person field organization. Visible early win. Substrate work needed is bounded.

Pod 1B Predictive Consumables Forecasting

OUTCOME	Reduce variance in customer consumables orders through AI-driven forecasting.
VALIDATE CRITERION	Standing-order variance for the top 50 customer accounts drops by 25% in the pilot quarter (Q4 2026).
ANATOMY	– Pod lead: Senior account manager – 1 Aura engineer – 1 inventory / logistics person – Periodic AI Office check-in
TIMELINE	8 weeks.
SUBSTRATE DEPENDENCIES	Customer consumables order history; production data where customers share it.
WHY FIRST WAVE	Direct working-capital impact. Stress-tests the customer-facing intelligence layer the strategy commits to.

Pod 2A — AI-Enabled Queueing Demo Build

OUTCOME	Ship a demo-ready AI-enabled queueing system to compete with Q-MATIC and Wavetec — predictive routing, Arabic voice agent, natural-language manager analytics.
VALIDATE CRITERION	A demonstrable AI-enabled queueing demo presented in 5 customer tenders by week 12, with measurable improvement in tender close rate vs. baseline.
ANATOMY	– Pod lead: ES division product person – 2 Aura engineers – 1 senior salesperson (banking vertical primary) – 1 banking-vertical SME
TIMELINE	12 weeks.
SUBSTRATE DEPENDENCIES	WAM platform integration access; Arabic NLP capability (vendor or build); customer demo environment.
WHY FIRST WAVE	This is the load-bearing pod for the AI Products workstream and the visible answer to the competitive scan. Its success or failure shapes the Phase 2 product roadmap heavily.

Pod 2B Sales Proposal Generation Engine

OUTCOME	Cut proposal drafting time for ES tenders by 90%+.
VALIDATE CRITERION	Average proposal generation time drops from approximately 2 days to under 1 hour by week 8, with quality maintained per sales lead review.
ANATOMY	– Pod lead: Senior sales rep – 1 Aura engineer – 1 ES product person – 1 sales operations lead
TIMELINE	6–8 weeks (shortest pod; fastest validate cycle).
SUBSTRATE DEPENDENCIES	Past proposals; product configuration data; pricing/discount guidelines.
WHY FIRST WAVE	Quick win. Demonstrates AI capability immediately to a non-engineering function (sales). Generates the kind of organic pull the strategy depends on.

Pod 3A — Sales & Customer Service Knowledge Layer

OUTCOME	Reduce time-to-answer in customer service and sales operations through AI-mediated knowledge access.
VALIDATE CRITERION	Customer service tier-1 response time drops by 40%, and sales rep time spent on internal lookups drops by 50%, by week 8.
ANATOMY	– Pod lead: Customer service lead – 1 Aura engineer – 1 sales operations person – 1 cross-divisional SME (one from each division)
TIMELINE	8 weeks.
SUBSTRATE DEPENDENCIES	CRM data; customer service ticket history; product documentation.
WHY FIRST WAVE	Broadest internal impact across functions. Provides an early reality-check on data accessibility constraints. Builds substrate that later operational pods will depend on.

Substrate Work

Substrate work runs in parallel with the pods, owned by the AI Office. It is prioritized by pod blockers, not by abstract architecture.

Substrate Stream 1: AI Development Environment

Weeks 1–3. Standing up the central AI development environment pods build on.

- Vendor/model access: foundation model APIs, vendor selection per use case
- Key management: secure API key storage and rotation
- Pod scaffolding: standard templates for kickoff, validate criteria, status reporting
- Logging and monitoring: shared observability for pod work in progress

Owned by the AI Office, supported by Aura engineering.

Substrate Stream 2: Data Access Layer

Ongoing, prioritized by pod blockers. Building the access layer pods need to ship.

- Dynamics AX read access: financial and operational data, exposed via API or SQL layer
- Customer service / CRM data: real-time access for Workstreams 1 and 3
- Field service ticket history: needed for Pod 1A
- Customer order history: needed for Pod 1B
- Product documentation: needed for Pods 2B and 3A

Owned by the AI Office in coordination with Aura engineering and the existing IT/operations function.

Substrate Stream 3: Arabic NLP Capability

Weeks 4–12. Needed for Pod 2A and subsequent customer-facing products.

- Vendor evaluation (Cohere Aya, Jais, custom Arabic-tuned models)
- Cost/quality benchmarking on Alfa-relevant use cases
- Selection and integration

Owned by Pod 2A in collaboration with the AI Office.

Substrate Stream 4: Shared Knowledge Base

Weeks 4–8. Capturing pod learnings so subsequent pods inherit them.

- Pod completion templates: what was tried, what worked, what did not
- Pattern library: emerging patterns across pods
- Cross-pod weekly sync notes

Owned by the AI Office.

Engineering Capacity Reality

A practical note worth flagging. With five pods running concurrently, each requiring 1–2 Aura engineers, Phase 1 will require approximately 7–9 engineers concurrently engaged from Aura. Given Aura's current size and existing commitments, this is a real constraint.

Mitigations:

- Pods staggered so no more than three pods are in their most engineering-intensive weeks simultaneously.
- Pods 2B and 3A are deliberately shorter and front-loaded so their engineering capacity frees up by mid-Phase 1.
- A clear policy: pods may NOT recruit additional Aura engineers mid-flight. Scope is constrained at kickoff. If a pod realizes it needs more, it goes to the Council for a scope decision — narrow scope, defer scope, or extend timeline.

If the engineering capacity reality forces a tradeoff between Phase 1 ambition and Phase 1 deliverability, the explicit guidance is: **deliver fewer pods well rather than more pods partially.** Three pods successfully validated is a stronger Phase 1 than six pods half-shipped.

30 / 60 / 90 / 180 Day Cadence

The Phase 1 rhythm has specific checkpoints.

Day 0 — Kickoff (Q3 2026 launch)

- Chairman opens with framing remarks.
- AI Council convenes formally for the first time; charter and authorities confirmed.
- Five Phase 1 pods named publicly with their pod leads.
- AI Office team named, two to three people initially.
- Champions Network: initial wave of champions announced.
- Substrate Stream 1 (AI Development Environment) begins.

Day 30

- All five pods have had their week-2 Council check-in.
- Substrate Stream 1 substantially complete; pods have working environments.
- Initial substrate findings reported (where data accessibility is harder than expected).
- First quarterly Sponsorship briefing prepared.

Day 60

- Mid-pod validates for Pods 2B and 3A (the fastest pods).
- Pod 1B mid-pod validate.
- Council makes first scale/stop/extend decisions.
- Substrate Stream 2 (data access layer) prioritized work substantially complete.
- Quarterly Sponsorship briefing delivered.

Day 90

- Pods 2B and 3A complete validate cycles. Council decides scale / merge / stop.
- Pod 1B complete validate cycle.
- Pods 1A and 2A enter their final validate weeks.
- Phase 1 mid-point review by the AI Council.
- First Champions Network gathering (lessons from the first wave).
- Wave 2 pods scoped, prepared for Q4 launch.

Day 120–180

- Pods 1A and 2A complete validate cycles.
- Phase 1 full retrospective by AI Council and Sponsorship layer.
- Calibration: which strategic assumptions held, which need revision.
- Phase 2 portfolio approved.
- First-half year strategic review delivered to the board level.

Early-Warning Indicators

Phase 1 success depends on catching problems early. Specific signals to watch for in weeks 2–4:

- **No pod has produced a demonstrable working draft by week 4.** Either pod scoping is too ambitious or substrate is more constrained than understood. *Action:* Council convenes early to re-scope.
- **More than 50% of pods are blocked on data accessibility.** Substrate priorities are wrong or substrate capacity is insufficient. *Action:* AI Office's substrate priorities reset.
- **The Champions Network is concentrated in engineering despite design intent.** Distribution effort needs reinforcement from sponsorship. *Action:* Chairman or Strategy Leader visibly engages non-engineering champion recruitment.
- **Aura engineering is consistently over-capacity.** Phase 1 ambition exceeds reality. *Action:* Council narrows scope per the "fewer pods well" principle.
- **One workstream is dominating Council attention.** Strategic risk materializing. *Action:* Council reviews workstream balance, adjusts.

Phase 1 Success Criteria

Phase 1 is deemed successful if, by year-end 2026:

1. At least three of the five proposed pods have completed their validate cycle with passing criteria.
2. At least one customer-facing AI capability has shipped to pilot or production — most likely the queueing AI demo from Pod 2A or the sales proposal engine from Pod 2B.
3. The substrate work has produced a foundation subsequent pods can build on without each pod re-inventing access.
4. The AI Council is operating as a real governance body, making scale/stop decisions on validate-criteria results rather than on political negotiation.

5. The Champions Network is real: at least one named, active champion in each major function across both divisions.
6. R&D investment trajectory is visible: spend in Phase 1 exceeds 2.16% of prorated revenue, on a path toward the multi-year target.

Phase 1 is **not** expected to deliver:

- Revenue growth in ES (that is a Phase 2 signal).
- Major Markem-Imaje partnership commitments (Phase 2 signal).
- Company-wide adoption (Phase 3+).
- A perfectly validated R&D investment ratio (takes 24 months).

The point of Phase 1 is to make the strategy real enough that Phase 2 can be planned from evidence rather than from theory.

SOURCES

References

This section documents the external sources behind the factual claims in this strategy — competitor AI capabilities, industry benchmarks, and the analytical frameworks the approach draws on. Internal claims (Alfa's financials, headcount, product mix, and the data-accessibility findings) are drawn from the company's own FY2025 records and the current-state assessment conducted for this engagement, and are not external citations.

Competitor AI Capabilities

The competitive scan reflects publicly available product information and announcements as of June 2026. Full detail is in the Competitor AI Scan appendix.

QUEUEING & CUSTOMER-FLOW VENDORS

- **Qmatic** — Aiva AI Voice Agent and the Qmatic Experience Cloud platform. Company product pages and press materials, 2025–2026.
- **Wavetec** — AI-powered Forecasting and Simulation module (announced May 2026), the Waillo conversational agent, and Waillo Insights natural-language analytics. Company announcements, 2026.
- **Tensator** — partnership with MODI Vision for facial-recognition-based passenger flow (announced April 2025); Virtual Assistant within the VQMS platform. Company announcements, 2025.
- **NEMO-Q** — virtual queuing and notification capabilities. Company product materials.

INDUSTRIAL CODING & MARKING VENDORS

- **Domino Printing Sciences** — R-Series machine vision inspection systems (following the July 2020 acquisition of Lake Image Systems), with the 1,000th unit milestone reported; Gx-Series PRO

printhead for the GS1 Digital Link / 2D code transition (launched April 2026, ahead of interpack 2026). Company announcements and trade-press coverage, 2020–2026.

- **Markem-Imaje** — AI inspection software (showcased at Empack 2026, April 2026); 24/7 AI-powered virtual assistant (launched October 2024); connected-printer IoT services; the CoLOS® software platform; and GS1 Digital Link positioning ahead of the 2027 sunrise. Company materials and trade-press coverage, 2024–2026.

Note: brand and product names are the trademarks of their respective owners and are referenced here for competitive-analysis purposes only.

Industry Benchmarks

R&D investment intensity — software-led B2B vs. distribution. Independent benchmarking surveys of private B2B SaaS companies place median R&D investment well above the level a hardware-distribution business typically carries:

- SaaS Capital's annual spending survey of 1,000+ private B2B SaaS companies (15th edition, March 2026) reports median R&D spend of roughly 22% of revenue, consistently the single largest expense category regardless of company size or funding model.
- BenchMarkit's 2024 B2B SaaS Performance Metrics report (also ~1,000 companies) reports a median nearer 31%.
- Analyses of SaaS companies at IPO (Blossom Street Ventures; Sammy Abdullah) converge on a median around 24% of revenue, with smaller firms spending materially more.

By contrast, hardware-centric and distribution businesses run far lower R&D intensity, with the economics driven by cost of goods rather than product development: industry margin analyses place hardware gross margins in the 20–40% range against software/SaaS gross margins of 60–90%.

Implication for the strategy's framing. The strategy's working figure — a step-up from Alfa's current 2.16% of revenue toward a software-led range — is deliberately conservative. The benchmark data would support a higher target. The 6–10% range used in this document represents a pragmatic 24-month trajectory for a company beginning from a distribution cost structure, not the eventual destination a fully software-led identity would imply. It is set to be credible and fundable from the existing recurring base, not to match pure-play SaaS peers in the first phase.

Benchmark sources: SaaS Capital 2026 Spending Benchmarks; BenchMarkit 2024 B2B SaaS Performance Metrics; Blossom Street Ventures and related IPO-cohort analyses; technology-sector gross-margin analyses, 2025–2026. Figures are medians from independent surveys and vary by company size, growth stage, and funding model.

Regulatory & Industry Timeline

- **GS1 Digital Link / 2D code transition ("2027 sunrise").** The industry-wide migration toward 2D codes (such as QR codes carrying GS1 Digital Link) for retail and supply-chain identification, with 2027 as the widely cited target milestone. Referenced throughout the strategy as a near-term forcing function for the coding and marking market. Source: GS1 standards organization materials and industry coverage.

Analytical Frameworks

The strategy's design draws on established frameworks in organizational change, innovation diffusion, adult learning, and AI capability:

- **Kotter's 8-Step Change Model** — John P. Kotter, *Leading Change* (Harvard Business Review Press). Informs the overall change architecture.
- **Diffusion of Innovations** — Everett M. Rogers, *Diffusion of Innovations* (Free Press). Informs the adopter-segmentation logic behind the Champions Network and pull-based adoption.
- **Andragogy (adult learning theory)** — Malcolm S. Knowles, *The Adult Learner*. Informs the learning-by-doing approach over formal curriculum.
- **The 70-20-10 model of learning and development** — originating from research associated with the Center for Creative Leadership (Lombardo & Eichinger). Informs the experiential/social/formal balance of the capability-building approach.
- **AI Fluency framework — Delegation, Description, Discernment, Diligence** — Anthropic's framework for effective individual AI use. Informs the section on what skilled AI use looks like.
- **AI maturity models** — Gartner and MIT Sloan Management Review maturity-model concepts (Awareness → Active → Operational → Systemic → Transformational). Informs the assessment of departmental AI maturity.

A Note on Method

Where this document describes patterns from "comparable organizations" and enterprise AI adoption efforts (in the Evidence and Reasoning section), it draws on the general body of public reporting on enterprise AI programs from 2023 to 2026 rather than on any single proprietary study. Specific company outcomes are not named, because the intent is to reflect the stabilized patterns across many public attempts — both successes and failures — rather than to rest the argument on any one case that might not transfer to Alfa's context.